

Exercise 1 Trees and heaps

This exercise is optional. Its purpose is to recall some facts about basic data structures, and to think about their properties.

- (a) (0 pts) Recall how a *binary heap* works, including the implementations of *insert*, *extract-min*, and building a heap in linear time. Recall how a *binary search tree* works, including the operations *search*, *insert*, and *delete*. Recall some dynamically balanced binary search tree, such as AVL- or red-black tree.
- (b) (2 extra pts) Let us define the *depth* of a node x in a binary search tree as the number of edges on the path from x to the root of the tree. Define the *subtree-size* of a node x to be the number of nodes in the subtree rooted at x (not counting x itself). The subtree-size of a leaf is thus 0 and the subtree-size of the root is one less than the number of nodes in the tree. Show that in *any* binary search tree, the sum of node-depths equals the sum of subtree-sizes.
Hint: Double-counting.
- (c) (2 extra pts) Recall the *rotation* operation in binary search trees. Show that given two arbitrary binary search trees T_1 and T_2 with nodes $\{1, \dots, n\}$, there is a sequence of at most $2n$ rotations that transforms T_1 into T_2 .
Hint: Rotate to some canonical state.
- (d) (2 extra pts) A *treap* is a binary tree in which every node stores a *pair* of values. The nodes of the treap satisfy the binary search tree order with respect to the first value of each pair, and they satisfy the (min)heap-order with respect to the second value of each pair. Construct a treap with the following pairs of values: $(3, 5), (1, 4), (2, 8), (9, 1), (8, 3), (6, 2), (4, 7), (5, 9), (7, 6)$. Is the treap unique?

Exercise 2 Stacks and variants

8 Points

Recall the array-based stack implementation from the lecture, with operations *create_stack*, *push*, *pop*, supported in constant amortized time.

- (a) (4 pts) A *double-ended queue* (“deque”) is a more powerful data structure that allows adding and removing elements at both ends of a list, i.e. it supports the operations *make-deque*, *push-right*, *push-left*, *pop-right*, *pop-left*. Design a deque that uses memory efficiently and in which all operations take constant amortized time. Analyze the data structure rigorously with any method you prefer. You may re-use the stack construction from the lecture as a black box, or you may use arrays.

- (b) (4 pts) Suppose you wanted a stack where all operations take *actual* constant time (as opposed to amortized). Describe how to modify the array-based design to achieve this. (A linked list would also work, but we insist on using arrays.)

Hint: The only “bad case” was when the array had to be doubled and elements had to be copied. Can you “spread” this work across multiple operations?

- (c) (3 extra pts) The *waste* of a data structure is the difference between the number of memory cells in use and the number n of items stored in the data structure. The stack implementation discussed in the lecture has a waste of $O(n)$. Improve the design to reduce the waste to $O(\sqrt{n})$ at all times. (Make sure you account for all the extra pointers and other bookkeeping you may need in the design.)

Exercise 3 Heap variants

12 pts

- (a) (6 pts) Design an efficient data structure that allows inserting keys, and the operations *extract-small*, and *extract-large*, which extract (i.e. output and remove from the data structure) an arbitrary key that is among the smallest 25%, respectively largest 25% of the keys currently in the data structure. Show that all operations take constant amortized time.

Hint: using a few stacks and a few additional memory cells may be sufficient.

- (b) (6 pts) Design an efficient data structure that supports the operations *insert* and *extract-median*, i.e. inserting a key, and returning and deleting the median of the current keys. You may use heaps as a black box. What are the running times of operations?

Exercise 4 Properties of heaps

6 Points

- (a) (2 pts) Recall that (in the comparison model) if one of the operations *insert* and *extract-min* take time $o(\log n)$, the other type of operation must take time $\Omega(\log n)$.

Similarly argue that if *meld* takes time $O(n^{1-\varepsilon})$ for arbitrary constant $\varepsilon > 0$, then *extract-min* must take time $\Omega(\log n)$.

- (b) (4 pts) Show the four properties of the binomial tree B_k observed in the lecture.

Hint: Induction.

From: Yumeng Li and Thore Brehmer

- 1) (b) (2 extra pts) Let us define the *depth* of a node x in a binary search tree as the number of edges on the path from x to the root of the tree. Define the *subtree-size* of a node x to be the number of nodes in the subtree rooted at x (not counting x itself). The subtree-size of a leaf is thus 0 and the subtree-size of the root is one less than the number of nodes in the tree. Show that in *any* binary search tree, the sum of node-depths equals the sum of subtree-sizes.

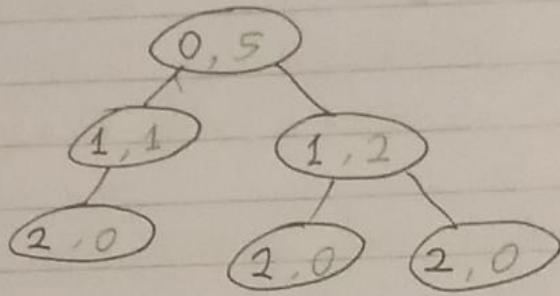
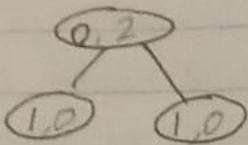
Hint: Double-counting.

- (c) (2 extra pts) Recall the *rotation* operation in binary search trees. Show that given two arbitrary binary search trees T_1 and T_2 with nodes $\{1, \dots, n\}$, there is a sequence of at most $2n$ rotations that transforms T_1 into T_2 .

Hint: Rotate to some canonical state.

Ex1 b) Depth Subtree-size

do a small example



A node x with depth i contributes to the total sum of depths by i . x has i ancestors $\Rightarrow x$ contributes to the total sum of subtree-sizes by i (since it contributes by 1 to the subtree-size of each of its ancestors)

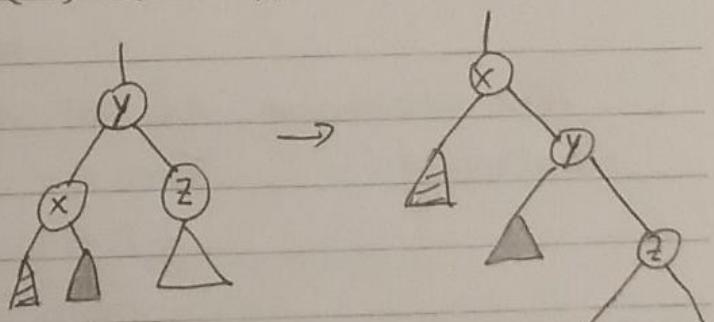
$$\text{total sum of depths} = \sum_{x \in \text{tree}} \text{depth}(x) = \sum_{x \in \text{tree}} \begin{cases} x \text{ contribution} \\ \text{to the total subtree-sizes sum} \end{cases}$$
$$= \text{total Subtree sizes sum.}$$

c). Idea : Transforming any tree of size n to a right spine chain can be done in $\leq n$ rotations by followings:

Algo: [When there exists a left branch from y to x , perform a right rotation at y .]

Note : After each iteration, we have one more right edge and one less left edge.

Observe : a tree w/ n nodes has $n-1$ edges. So we are done within ~~n steps~~ $(n-1)$ iterations.

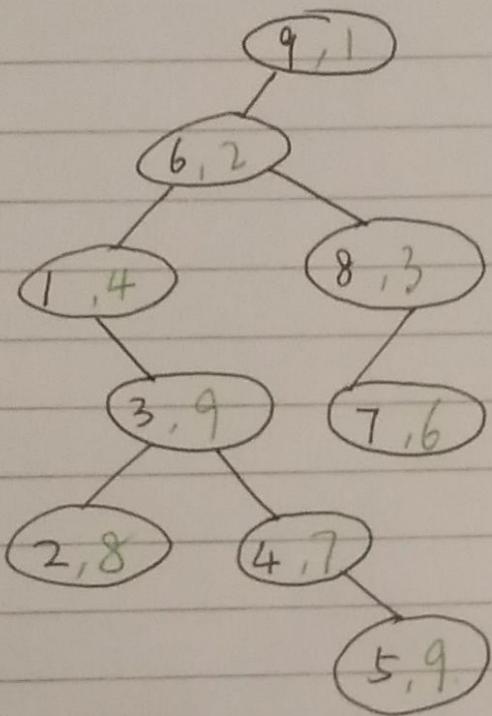


Algo.
 $T_1 \xrightarrow{(n-1) \text{ steps}} \text{right spine} \xrightarrow{(n-1) \text{ step}} T_2$

$\Rightarrow \leq 2(n-1)$ rotations to transform from T_1 to T_2 .

- (d) (2 extra pts) A *treap* is a binary tree in which every node stores a *pair* of values. The nodes of the treap satisfy the binary search tree order with respect to the first value of each pair, and they satisfy the (min)heap-order with respect to the second value of each pair. Construct a treap with the following pairs of values: (3, 5), (1, 4), (2, 8), (9, 1), (8, 3), (6, 2), (4, 7), (5, 9), (7, 6). Is the treap unique?

d)



The treap is unique

Recall the array-based stack implementation from the lecture, with operations `create_stack`, `push`, `pop`, supported in constant amortized time.

- (a) (4 pts) A *double-ended queue* ("deque") is a more powerful data structure that allows adding and removing elements at both ends of a list, i.e. it supports the operations `make_deque`, `push-right`, `push-left`, `pop-right`, `pop-left`.

Design a deque that uses memory efficiently and in which all operations take constant amortized time. Analyze the data structure rigorously with any method you prefer. You may re-use the stack construction from the lecture as a black box, or you may use arrays.

Operations:

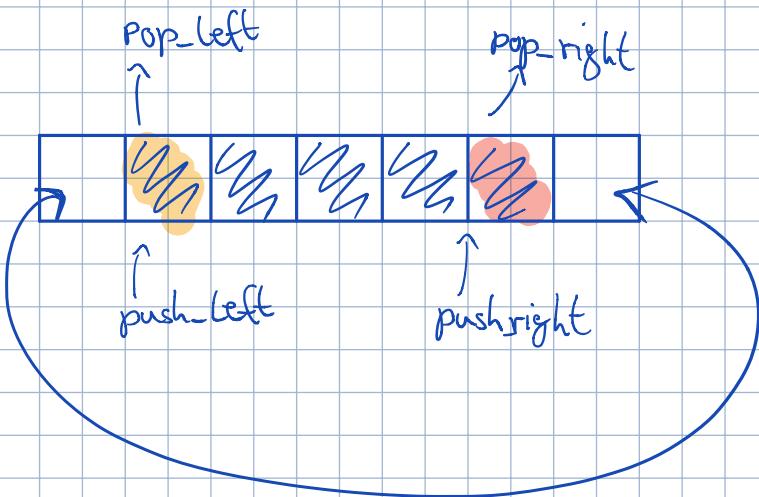
(x1) `make_deque`

`push-right`

`push-left`

`pop-right`

`pop-left`



Pointers:

`left of stack` (as int → initially 0)

`right of stack` (as int → initially 0)

`capacity` = size of array → initially 3

`n` = number of elements → initially 0

Invariant

$$\lfloor \frac{\text{capacity}}{4} \rfloor \leq n \leq \text{capacity}$$

def pop-left(A)

if $n = 0$:

O(1) → nothing to pop.
Stack is empty

else:

if $\text{left} = \text{capacity} - 1$:

O(1) return $A[\text{left}]$

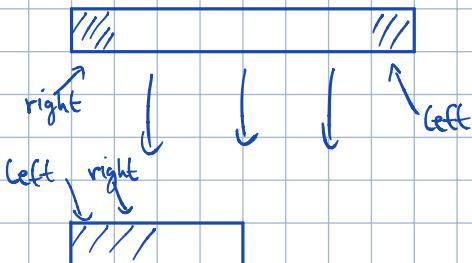
$\text{left} = 0$

else:

O(1) return $A[\text{left}]$

$\text{left} += 1$

if $\lfloor \frac{\text{capacity}}{4} \rfloor = n$:



allocate array of size $\lfloor \frac{n}{2} \rfloor$

and copy all elements over in order starting from

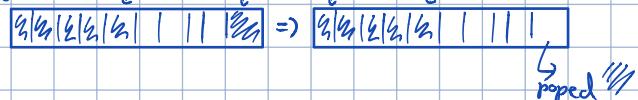
(left pointer) to right (pointer)

(pop-right is really similar to pop-left)

Obs: Then: 
To get to this case

Now: 

e.g. right left left right



left right



Accounting Method

→ pop will cost 3ϵ

↳ 1ϵ for simple operation (popping)

↳ 2ϵ stored in item for later

$\Rightarrow \frac{n}{2}$ push operations will store $2\epsilon \cdot \frac{n}{2} = n\epsilon$ total in the items

- To allocate a $\frac{capacity}{2}$ long list. There will be n copies $\Rightarrow n\epsilon$ cost

\Rightarrow we have exactly as much ϵ as we need $\Rightarrow O(1)$

def push-left(v, A)

if (left == 0):

$O(n)$

$$L_{\text{eff}} = \text{capacity} - 1$$

$A[\text{left}] = \text{left}$

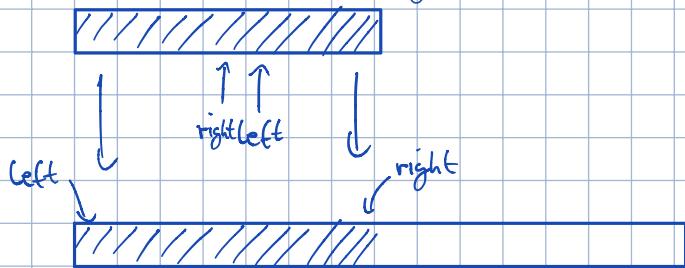
else :

$e(n)$

left = 1

`4 [left] = Left`

if ($n == \text{capacity}$):



allocate array of size $2n$

and copy all elements over in order starting from
 $\text{left}(\text{pointer})$ to right (pointer)

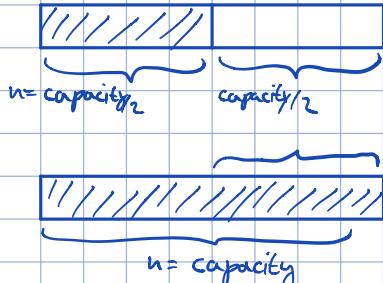
(push-right is just like push-left with minor differences)

Obs. There have been $3^{\lfloor \frac{n}{2} \rfloor}$ push (left or right) operations, before the array needs to be allocated to $2n$.

Then :

To get to this case

Now:



$\frac{n}{2}$ push-operations

Accounting Method

push will cost 3ϵ

↳ 1ϵ for simple operation (pushing)

↳ 2ϵ stored in item for later

$\Rightarrow \frac{n}{2}$ push operations will store $2\epsilon \cdot \frac{n}{2} = n\epsilon$ total in the items

- To allocate a 2 capacity long list. There will be n copies $\Rightarrow n\epsilon$ cost

\Rightarrow we have exactly as much ϵ as we need $\Rightarrow O(n)$

- (b) (4 pts) Suppose you wanted a stack where all operations take *actual* constant time (as opposed to amortized). Describe how to modify the array-based design to achieve this. (A linked list would also work, but we insist on using arrays.)

Hint: The only “bad case” was when the array had to be doubled and elements had to be copied. Can you “spread” this work across multiple operations?

Operations

$O(1)$ $t = \text{make_stack}()$

$\text{pop}(t)$

$\text{push}(v, t)$

Idea: when allocating a $2n$ long list, we don't copy all the elements from the full list to it at once.

We only copy one element from the full list to the $2n$ list, when we try to push a new element in it.

After we pushed n new elements to the $2n$ list, all n elements from the old (full) list have been copied as well.

- For that we need a pointer on the full list.

→ Which starts at $\text{old-list}[0]$

After the pointer reached the end of the old-list → The full list can be deleted

- We also need a pointer on the new $2n$ list, which shows where to push new elements to.

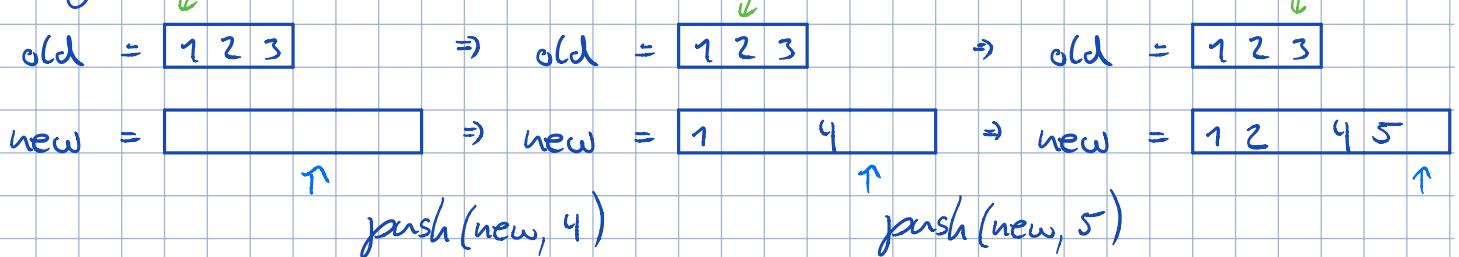
→ Starts at $\text{new-list}[n]$

After the pointer reached the end of the new-list → $\text{old-list} = \text{new-list}$ and $\text{new-list} = \text{allocate array with size } 2n$

- We need to change the push operation, so that it will copy the element of the old-list pointer ⇒ $O(1)$ (remove the case in which it normally copies all n elements)

- Need to add a case in pop , if we try to pop items which haven't been copied yet. Decrement the new-list pointer ⇒ still $O(1)$

e.g. 1



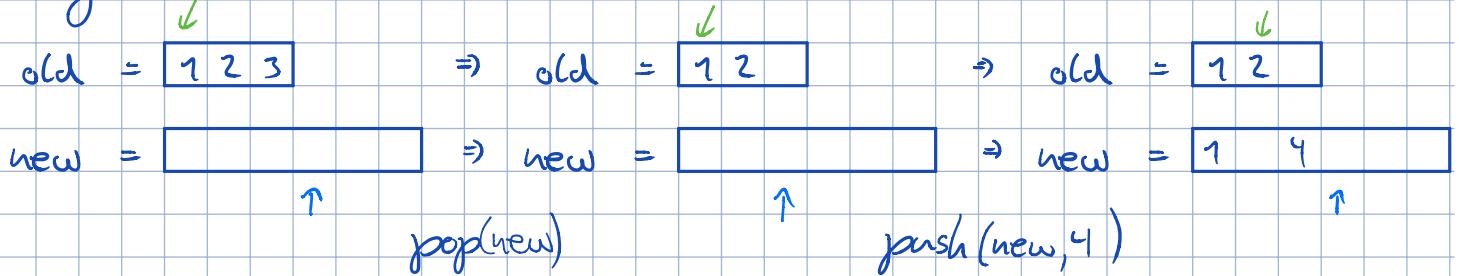
\Rightarrow ~~1 2 3~~ deleted

\Rightarrow old = $\boxed{1 \ 2 \ 3 \ 4 \ 5 \ 6}$

new = $\boxed{}$

push(new, 6) \uparrow

e.g. 2



\Rightarrow old = $\boxed{1 \ 2}$

\Rightarrow old = $\boxed{1 \ 2 \ 3}$

\Rightarrow new = $\boxed{1 }$

\Rightarrow new = $\boxed{1 \ 2 \ 5}$

pop(new)

push(new, 5)

- (c) (3 extra pts) The *waste* of a data structure is the difference between the number of memory cells in use and the number n of items stored in the data structure. The stack implementation discussed in the lecture has a waste of $O(n)$. Improve the design to reduce the waste to $O(\sqrt{n})$ at all times. (Make sure you account for all the extra pointers and other bookkeeping you may need in the design.)

Idea:

- Use *linked lists* instead of arrays.

(Increase List)

- instead of resizing the list to a list of size $2n$, we will add one more linked list of size $\lceil \frac{n}{2} \rceil - 1$. (-1 because the pointer of a linked list takes 1 space)

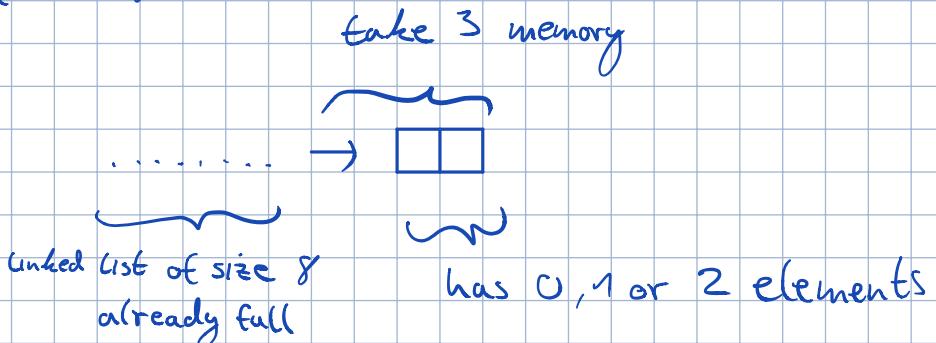
(Decrease List)

- Just delete the latest empty linked list.

Works for $n > 8$

E.g.

for $n = 8 \dots 10$



⇒ waste will be 3, 1 or 0

0 if both cells are filled

1 if one cell is filled

3 if both cells are empty

↳ (at this point pointer will count as waste as well)

$\sqrt{8} = 3$ waste is allowed.

- (a) (6 pts) Design an efficient data structure that allows inserting keys, and the operations *extract-small*, and *extract-large*, which extract (i.e. output and remove from the data structure) an arbitrary key that is among the smallest 25%, respectively largest 25% of the keys currently in the data structure. Show that all operations take constant amortized time.

3a: We use three stacks S, M, L with invariants:

$$|S| \leq 0.25n, |S| > 0.$$

n := current number of elements

$$|L| \leq 0.25n, |L| > 0.$$

$S \leq M \leq L$, (i.e. elements in S are smaller than ^{those} in M and elements in M are smaller than those in L)

And store "splitting" elements a, b st. $S \leq a, L \geq b$

(i.e. all elements in S are smaller than a , and all elements in L are smaller than b).

And keep n up to date.

Implementation: ① $\text{insert}(x)$: Compare x to a , if $x \leq a$, push x to S
 $O(1)$ + Time (Rebalance). if $x \geq b$, push x to L
else, push x to M .

$$n \leftarrow n+1$$

If $|S| > 0.25n$ or $|L| > 0.25n$, rebalance.

② extract-small : pop from S

$$n \leftarrow n-1$$

If $|S|=0$ or $|L| > 0.25n$, rebalance.

③ extract-large : symmetric to ②.

④ rebalance : Pop all elements in one list

time $O(n)$

Find elements a, b at rank $\frac{1}{8}n$ and $\frac{7}{8}n$.

Put elements $\leq a$ into S

Put elements $\geq b$ into L .

rest into M .

After rebalancing, we can insert at least m elements until next rebalancing. Then we should have

$$\frac{\frac{1}{8}n+m}{n+m} \leq \frac{1}{4}(n+m).$$

consider $m = \frac{1}{8}n$, this holds, so $m \geq \frac{1}{8}n$.

So Amortized time of rebalancing is $\frac{O(n)}{\frac{1}{8}n} = O(1)$.

- (b) (6 pts) Design an efficient data structure that supports the operations *insert* and *extract-median*, i.e. inserting a key, and returning and deleting the median of the current keys. You may use heaps as a black box. What are the running times of operations?

Idea: use a Min-heap that stores the items \leq median and use a Max-heap that stores the items \geq median

- $\text{extract-median}(\text{Max-heap}, \text{Min-heap}) \Rightarrow O(\log n)$

if $|\text{Max-heap}| > |\text{Min-heap}|$
 $O(1)$ return root of Max-heap
 $O(1)$ remove root of Max-heap
 $O(\log n)$ balance Max-heap

if $|\text{Max-heap}| < |\text{Min-heap}|$
 $O(1)$ return root of Min-heap
 $O(1)$ remove root of Min-heap
 $O(\log n)$ balance Min-heap

if $|\text{Max-heap}| = |\text{Min-heap}|$
 $O(1)$ return root of Max-heap
 $O(1)$ remove root of Max-heap
 $O(\log n)$ balance Max-heap

(\Rightarrow elements are even, there are two medians, doesn't matter if we take it from Max- or Min-heap)
 we just take it from Max-heap

- $\text{new-insert}(x, \text{Max-heap}, \text{Min-heap}) \Rightarrow O(\log n)$

if $|\text{Max-heap}| = |\text{Min-heap}| = 0$: (if both are empty just insert it in either one. We just take the Max-heap)
 $O(\log n)$ insert(x) in Max-heap

if $x \leq$ root of Max-heap:
 $O(\log n)$ insert(x) in Max-heap

if $x >$ root of Max-heap:
 $O(\log n)$ insert(x) in Min-heap

if $||\text{Max-heap}| - |\text{Min-heap}|| = 2$:

(If the size of Max- and Min-heap differ by 2. Take the root of the bigger heap and insert it in the smaller heap)

if $|\text{Max-heap}| > |\text{Min-heap}|$:
 $O(\log n)$ insert(root of Max-heap) in Min-heap
 $O(\log n)$ balance Max-heap

if $|\text{Max-heap}| < |\text{Min-heap}|$:
 $O(\log n)$ insert(root of Min-heap) in Max-heap
 $O(\log n)$ balance Min-heap

e.g. [2, 3, 7, 5, 4, 1, 6]

1 2 3 4 5 6 7

Max. Ht
insert(2) => (2)

Min. Ht

Max. Ht
insert(2) => (2)

Min. Ht
(3)

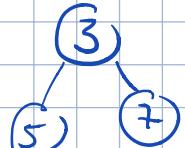
Max. Ht
insert(7) => (2)

Min. Ht



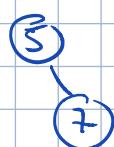
Max. Ht
insert(5) => (2)

Min. Ht
(3)



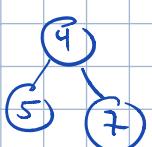
Max. Ht
=> (3)

Min. Ht



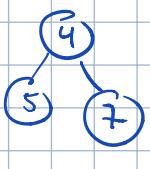
Max. Ht
insert(4) => (2)

Min. Ht
(3)



Max. Ht
insert(1) => (3)

Min. Ht



Max. Ht
insert(6) =>

Min. Ht
(4)



- (a) (2 pts) Recall that (in the comparison model) if one of the operations *insert* and *extract-min* take time $o(\log n)$, the other type of operation must take time $\Omega(\log n)$.

Similarly argue that if *meld* takes time $O(n^{1-\varepsilon})$ for arbitrary constant $\varepsilon > 0$, then *extract-min* must take time $\Omega(\log n)$.

Ex4 a). Recall = insert + extract-min \Rightarrow sorting.

Suppose we have n elements. We meld every two elements into a heap, each heap takes ~~$O(n^{1-\varepsilon})$~~ $O(1)$ time, and there are intotal $\frac{n}{2}$ heaps.

Then we meld these two-element heaps, each takes $O(2^{1-\varepsilon})$ time every two of and intotal $\frac{n}{4}$ heaps.

Do this iteration until we have an n -heap.

$$\begin{aligned} \text{Time : } & \frac{n}{2}O(1) + \frac{n}{4}O(2^{1-\varepsilon}) + \frac{n}{8}O(2^2)^{1-\varepsilon} + \dots \\ & = o\left(\frac{n}{2} \lg n\right). \end{aligned}$$

Since we have the heap, we can extract-min for n times

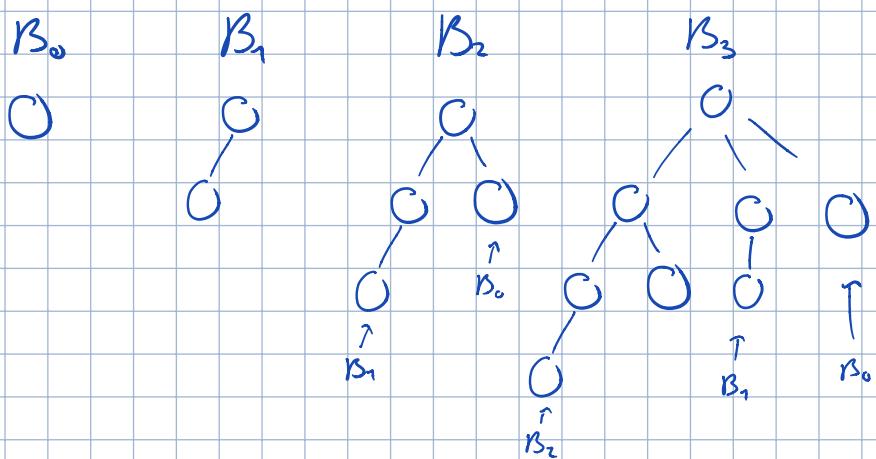
Then we ~~get~~ sort the n elements.

$$\text{So } o\left(\frac{n}{2} \lg n\right) + T(\text{extract-min}) \cdot n \geq O(\lg n \cdot n).$$

$$\text{So } T(\text{extract-min}) = \Omega(\lg n).$$

(b) (4 pts) Show the four properties of the binomial tree B_k observed in the lecture.

Hint: Induction.



Obs.: B_k has $\cdot 2^k$ nodes,
 • root of degree k ,
 • height k .
 $\binom{k}{e}$ nodes at depth e

Prof (Exercise)

BS_{k+1} : root node and children BS_{k-1}, \dots, BS_0 (in this order)

BS_k has 2^k nodes

J.A. $k=0$ $BS_0 = 2^0$ ✓

J.U. true for k

J.S. show for $k+1$

BS_{k+1} has 2^{k+1} nodes

= BS_{k+1} has $1 + [BS_k + BS_{k-1} + \dots + BS_0]$ nodes

$$= BS_k + \underbrace{1 + [BS_{k-1} + \dots + BS_0]}$$

$$= BS_k + BS_k$$

$$\text{J.U. } 2^k + 2^k = 2 \cdot 2^k = 2^{k+1} \quad \square$$

BS_0
 $\textcircled{1} \leftarrow 1 \text{ Node}$

B_n has root of degree k (don't really need induction for this)

J.A. $k=0$ B_0 has root of degree 0 ✓

J.U. True for k

J.S. show for $k+1$

B_{k+1} has root of degree $k+1$

= B_{k+1} has root of degree sum of B_k s $\underbrace{[B_k, B_{k-1}, \dots, B_0]}_{\text{any}}$

B_{k+1} has root of degree $1 + \sum_{i=0}^k 1$

B_{k+1} has root of degree $1+k$

□

B_n has height k

J.A. $k=0$ B_0 has height 0 ✓

B_0 height 0

J.U. True for k

J.S. show for $k+1$

B_{k+1} has height $k+1$

= B_{k+1} has height $1 + \max_{\text{height}} [B_k, B_{k-1}, \dots, B_0]$

J.U. = B_{k+1} has height $1 + \max_{\text{height}} [k + k-1 + \dots + 0]$

B_{k+1} has height $1+k$

□

B_k has $\binom{k}{l}$ nodes at depth k

J.A B_0 has $\binom{0}{c}$ nodes at depth 0
= 1

B_0

C) 1 node at depth 0

J.U. true for k

J.S. show for $k+1$

B_{k+1} has $\binom{k+1}{l}$ nodes at depth $k+1$

$B_{k+1} = \text{root}' + [B_2, B_{k+1}, \dots, B_0]$ ('+' meaning as connection with the roots)

$B_{k+1} = B_k (\text{but one 'level' lower})' + \text{root}' + [B_{k+1}, \dots, B_0]$

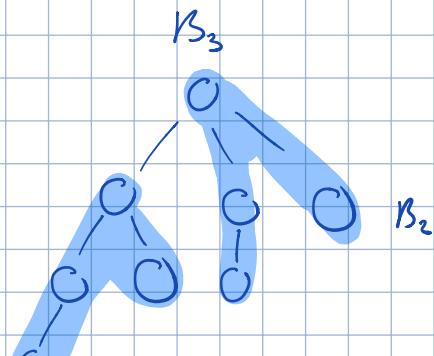
$B_{k+1} = B_k (\text{but one 'level' lower})' + B_k (-1 \text{ level as } B_{k+1}) + B_k (\text{same level as } B_{k+1})$

$\rightarrow J.U. = \binom{k}{k-1} + \binom{k}{k}$

$$= \binom{k+1}{k}$$

$$\left(\text{rule: } \binom{n+1}{k+1} = \binom{n}{k} + \binom{n}{k+1} \right)$$

e.g. for 'Level'



B_2 but one level lower than B_3